

## Problem 4, Parts C-E: Stochastic Gradient Descent Visualization

In this Jupyter notebook, we visualize how SGD works. This visualization corresponds to parts C-E of question 4 in set 1.

Use this notebook to write your code for problem 4 parts C-E by filling in the sections marked `# TODO` and running all cells.

```
In [81]: # Setup.

import numpy as np
import matplotlib.pyplot as plt
from IPython.display import HTML
import math

from sgd_helper import (
    generate_dataset1,
    generate_dataset2,
    plot_dataset,
    plot_loss_function,
    animate_convergence,
    animate_sgd_suite
)
```

### Problem 4C: Implementation of SGD

Fill in the loss, gradient, and SGD functions according to the guidelines given in the problem statement in order to perform SGD.

```

In [82]: def loss(X, Y, w):
    """
    Calculate the squared loss function.

    Inputs:
        X: A (N, D) shaped numpy array containing the data points.
        Y: A (N, ) shaped numpy array containing the (float) labels of the data points.
        w: A (D, ) shaped numpy array containing the weight vector.

    Outputs:
        The loss evaluated with respect to X, Y, and w.
    """
    predict = []
    for x in X:
        predict.append(np.inner(w, x))
    predict = np.asarray(predict)

    loss = 0
    for i in range(len(predict)):
        loss += (predict[i] - Y[i]) ** 2

    return loss

def gradient(x, y, w):
    """
    Calculate the gradient of the loss function with respect to
    a single point (x, y), and using weight vector w.

    Inputs:
        x: A (D, ) shaped numpy array containing a single data point.
        y: The float label for the data point.
        w: A (D, ) shaped numpy array containing the weight vector.

    Output:
        The gradient of the loss with respect to x, y, and w.
    """
    #=====
    # TODO: Implement the gradient of the loss function.
    #=====
    grad = -2 * (y - np.inner(w, x)) * x
    return grad

```

```

In [ ]: def SGD(X, Y, w_start, eta, N_epochs):
        """
        Perform SGD using dataset (X, Y), initial weight vector w_start,
        learning rate eta, and N_epochs epochs.

        Outputs:
            w: A (D, ) shaped array containing the final weight vector.
            losses: A (N_epochs, ) shaped array containing the losses from all iterations.
        """

        #=====
        # TODO: Implement the SGD algorithm.
        #=====

        totalLoss = []
        weights = w_start

        #start loss func (for some reason it won't work when i call the function?)
        #Python TypeError: 'list' object is not callable.
        #and
        #TypeError: 'numpy.float64' object is not callable
        predict = []

        for x in X:
            predict.append(np.inner(weights, x))
        predict = np.asarray(predict)

        loss = 0
        for i in range(len(predict)):
            loss += (predict[i] - Y[i]) ** 2
        #end loss func

        #totalLoss.append(loss)

        for n in range(N_epochs):
            for i in range(len(X)):
                g = gradient(X[i], Y[i], weights)
                weights -= eta * g

            #start loss func
            predict = []
            for x in X:
                predict.append(np.inner(weights, x))
            predict = np.asarray(predict)

            loss = 0
            assert(len(predict) == len(Y))
            for i in range(len(predict)):
                loss += (predict[i] - Y[i]) ** 2

            #currLoss = loss(X, Y, weights)
            totalLoss.append(loss)
            #print(str(n) + ": " + str(loss) + ", " + str(weights))

        return np.asarray(weights), np.asarray(totalLoss)


```

## Problem 4D: Visualization

## Dataset


We'll start off by generating two simple 2-dimensional datasets. For simplicity we do not consider separate training and test sets.

```
In [83]: X1, Y1 = generate_dataset1()
         plot_dataset(X1, Y1)
```



```
Out[83]: (<Figure size 432x288 with 2 Axes>,
         <matplotlib.axes._subplots.AxesSubplot at 0x11708fa90>)
```

```
In [84]: X2, Y2 = generate_dataset2()
         plot_dataset(X2, Y2)
```



```
Out[84]: (<Figure size 432x288 with 2 Axes>,
         <matplotlib.axes._subplots.AxesSubplot at 0x1170e0750>)
```

## SGD from a single point



First, let's visualize SGD from a single starting point:

```
In [85]: # Parameters to feed the SGD.  
# <FR> changes the animation speed.  
params = ({'w_start': [0.01, 0.01], 'eta': 0.00001},)  
N_epochs = 1000  
FR = 20  
  
# Let's animate it!  
anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR)  
HTML(anim.to_html5_video())
```

Performing SGD with parameters {'w\_start': [0.01, 0.01], 'eta': 1e-05} ...

Animating...

Out[85]:



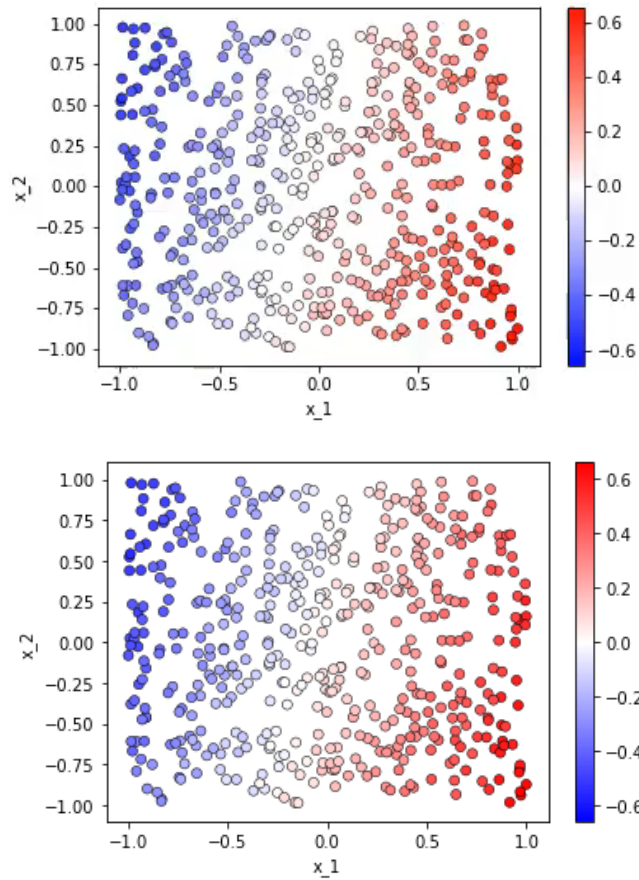
Let's view how the weights change as the algorithm converges:

```
In [86]: # Parameters to feed the SGD.
params = ({'w_start': [0.01, 0.01], 'eta': 0.00001},)
N_epochs = 1000
FR = 20

# Let's do it!
W, _ = SGD(X1, Y1, params[0]['w_start'], params[0]['eta'], N_epochs)
anim = animate_convergence(X1, Y1, W, FR)
HTML(anim.to_html5_video())
```

Animating...

Out[86]:



## SGD from multiple points

Now, let's visualize SGD from multiple arbitrary starting points:



```
In [87]: # Parameters to feed the SGD.
# Here, we specify each different set of initializations as a dictionary.
params = (
    {'w_start': [-0.8, -0.3], 'eta': 0.00001},
    {'w_start': [-0.9, 0.4], 'eta': 0.00001},
    {'w_start': [-0.4, 0.9], 'eta': 0.00001},
    {'w_start': [0.8, 0.8], 'eta': 0.00001},
)
N_epochs = 1000
FR = 20

# Let's go!
anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR)
HTML(anim.to_html5_video())
```

Performing SGD with parameters {'w\_start': [-0.8, -0.3], 'eta': 1e-05} ...  
 Performing SGD with parameters {'w\_start': [-0.9, 0.4], 'eta': 1e-05} ...  
 Performing SGD with parameters {'w\_start': [-0.4, 0.9], 'eta': 1e-05} ...  
 Performing SGD with parameters {'w\_start': [0.8, 0.8], 'eta': 1e-05} ...

Animating...

Out[87]:



Let's do the same thing but with a different dataset:



```
In [88]: # Parameters to feed the SGD.
params = (
    {'w_start': [-0.8, -0.3], 'eta': 0.00001},
    {'w_start': [-0.9, 0.4], 'eta': 0.00001},
    {'w_start': [-0.4, 0.9], 'eta': 0.00001},
    {'w_start': [0.8, 0.8], 'eta': 0.00001},
)
N_epochs = 1000
FR = 20

# Animate!
anim = animate_sgd_suite(SGD, loss, X2, Y2, params, N_epochs, FR)
HTML(anim.to_html5_video())

Performing SGD with parameters {'w_start': [-0.8, -0.3], 'eta': 1e-05} ...
Performing SGD with parameters {'w_start': [-0.9, 0.4], 'eta': 1e-05} ...
Performing SGD with parameters {'w_start': [-0.4, 0.9], 'eta': 1e-05} ...
Performing SGD with parameters {'w_start': [0.8, 0.8], 'eta': 1e-05} ...

Animating...
```

Out[88]:



## Problem 4E: SGD with different step sizes

Now, let's visualize SGD with different step sizes ( $\eta$ ):

(For ease of visualization: the trajectories are ordered from left to right by increasing  $\eta$  value. Also, note that we use smaller values of  $N_{\text{epochs}}$  and  $FR$  here for easier visualization.)





```
In [89]: # Parameters to feed the SGD.
params = (
    {'w_start': [0.7, 0.8], 'eta': 0.00001},
    {'w_start': [0.2, 0.8], 'eta': 0.00005},
    {'w_start': [-0.2, 0.7], 'eta': 0.0001},
    {'w_start': [-0.6, 0.6], 'eta': 0.0002},
)
N_epochs = 100
FR = 2

# Go!
anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR, ms=2)
HTML(anim.to_html5_video())

Performing SGD with parameters {'w_start': [0.7, 0.8], 'eta': 1e-05} ...
Performing SGD with parameters {'w_start': [0.2, 0.8], 'eta': 5e-05} ...
Performing SGD with parameters {'w_start': [-0.2, 0.7], 'eta': 0.0001} ...
Performing SGD with parameters {'w_start': [-0.6, 0.6], 'eta': 0.0002} ...

Animating...
```

Out[89]:



## Plotting SGD Convergence

Let's visualize the difference in convergence rates a different way. Plot the loss with respect to epoch (iteration) number for each value of eta on the same graph.

```

In [90]: '''Plotting SGD convergence'''

#=====
# TODO: For the given learning rates, plot the
# loss for each epoch.
#=====

eta_vals = [1e-6, 5e-6, 1e-5, 3e-5, 1e-4]
w_start = [0.01, 0.01]
N_epochs = 1000


weights = []
loss = []

for eta in eta_vals:
    initial = [0.001, 0.001, 0.001, 0.001, 0.001]
    finalWeights, allLoss = SGD(X1, Y1, w_start, eta, N_epochs)
    weights.append(finalWeights)
    loss.append(allLoss)

fig = plt.figure()
x = epochs
plt.title('eta vs Loss', fontsize = 22)
plt.plot(loss[0], marker = '.')
plt.plot(loss[1], marker = '.')
plt.plot(loss[2], marker = '.')
plt.plot(loss[3], marker = '.')
plt.plot(loss[4], marker = '.')
plt.legend(('e-6 rate', '5e-6 rate', 'e-5 rate', '3e-5 rate',
'e-4 rate'), loc = 'upper right', fontsize = 14)
plt.xscale('log')
plt.ylabel('Loss', fontsize = 18)

```

Out[90]: Text(0, 0.5, 'Loss')



Clearly, a big step size results in fast convergence! Why don't we just set eta to a really big value, then? Say, eta=1?

(Again, note that the FR is lower for this animation.)

```
In [91]: # Parameters to feed the SGD.
params = ({'w_start': [0.01, 0.01], 'eta': 1},)
N_epochs = 100
FR = 2

# Voila!
anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR, ms=2)
HTML(anim.to_html5_video())

Performing SGD with parameters {'w_start': [0.01, 0.01], 'eta': 1} ...

-----
TypeError                                Traceback (most recent call last)
<ipython-input-91-6f8ee3a8be68> in <module>
      5
      6 # Voila!
----> 7 anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR, ms=2)
      8 HTML(anim.to_html5_video())

~/Downloads/CS155_SET1/sgd_helper.py in animate_sgd_suite(SGD, loss, X, Y, param
s, N_epochs, FR, ms)
    121
    122     # Get the loss grid and plot it.
--> 123     w_grid, loss_grid = get_loss_grid((-1, 1, 100), (-1, 1, 100), X, Y,
loss)
    124     fig, ax = plot_loss_function(w_grid[0], w_grid[1], loss_grid)
    125

~/Downloads/CS155_SET1/sgd_helper.py in get_loss_grid(x_params, y_params, X, Y,
loss)
    77         for j in range(len(loss_grid[0])):
    78             w = np.array([w_grid[0][i, j], w_grid[1][i, j]])
----> 79             loss_grid[i, j] = loss(X, Y, w)
    80
    81     return w_grid, loss_grid

TypeError: 'list' object is not callable
```

Just for fun, let's try eta=10 as well. What happens? (Hint: look at W)

```
In [ ]: # Parameters to feed the SGD.
w_start = [0.01, 0.01]
eta = 10
N_epochs = 100

# Presto!
W, losses = SGD(X1, Y1, w_start, eta, N_epochs)
```

## Extra Visualization (not part of the homework problem)

One final visualization! What happens if the loss function has multiple optima?

```
In [ ]: # Import different SGD & loss functions.
        # In particular, the loss function has multiple optima.
        from sgd_multiopt_helper import SGD, loss

        # Parameters to feed the SGD.
        params = (
            {'w_start': [0.9, 0.9], 'eta': 0.01},
            {'w_start': [0.0, 0.0], 'eta': 0.01},
            {'w_start': [-0.8, 0.6], 'eta': 0.01},
            {'w_start': [-0.8, -0.6], 'eta': 0.01},
            {'w_start': [-0.4, -0.3], 'eta': 0.01},
        )
        N_epochs = 100
        FR = 2

        # One more time!
        anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR, ms=2)
        HTML(anim.to_html5_video())
```

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]: