

Problem 4, Parts F-H: Stochastic Gradient Descent with a Larger Dataset

Use this notebook to write your code for problem 4 parts F-H by filling in the sections marked `# TODO` and running all cells.

```
In [3]: # Setup.  
import csv  
import numpy as np  
import matplotlib.pyplot as plt  
import math  
import random  
  
%matplotlib inline
```

Problem 4F: Perform SGD with the new dataset

For the functions below, you may re-use your code from parts 4C-E. Note that you can now modify your SGD function to return the final weight vector instead of the weights after every epoch.

```

In [4]: def loss(X, Y, w):
        '''
        Calculate the squared loss function.

        Inputs:
            X: A (N, D) shaped numpy array containing the data points.
            Y: A (N, ) shaped numpy array containing the (float) labels of the data points.
            w: A (D, ) shaped numpy array containing the weight vector.

        Outputs:
            The loss evaluated with respect to X, Y, and w.
        '''
        predict = []
        for x in X:
            predict.append(np.inner(w, x))
        predict = np.asarray(predict)

        loss = 0
        for i in range(len(predict)):
            loss += (predict[i] - Y[i]) ** 2

        return loss

def gradient(x, y, w):
    '''
    Calculate the gradient of the loss function with respect to
    a single point (x, y), and using weight vector w.

    Inputs:
        x: A (D, ) shaped numpy array containing a single data point.
        y: The float label for the data point.
        w: A (D, ) shaped numpy array containing the weight vector.

    Output:
        The gradient of the loss with respect to x, y, and w.
    '''

    #=====
    # TODO: Implement the gradient of the loss function.
    #=====
    grad = -2 * (y - np.inner(w, x)) * x
    return grad

```

```

In [ ]: def SGD(X, Y, w_start, eta, N_epochs):
        '''
        Perform SGD using dataset (X, Y), initial weight vector w_start,
        learning rate eta, and N_epochs epochs.

        Outputs:
            w: A (D, ) shaped array containing the final weight vector.
            losses: A (N_epochs, ) shaped array containing the losses from all iterations.
        '''

        #=====
        # TODO: Implement the SGD algorithm.
        #=====

        totalLoss = []
        weights = w_start

        #start loss func (for some reason it won't work when i call the function?)
        #Python TypeError: 'list' object is not callable.
        #and
        #TypeError: 'numpy.float64' object is not callable
        predict = []
        for x in X:
            predict.append(np.inner(weights, x))
        predict = np.asarray(predict)

        loss = 0
        assert(len(predict) == len(Y))
        for i in range(len(predict)):
            loss += (predict[i] - Y[i]) ** 2
        #end loss func

        totalLoss.append(loss)

        for n in range(N_epochs):
            assert(len(X) == 1000)

            for i in range(len(X)):
                g = gradient(X[i], Y[i], weights)
                assert(len(g) == len(weights))
                weights -= eta * g

            #start loss func
            predict = []
            for x in X:
                predict.append(np.inner(weights, x))
            predict = np.asarray(predict)

            loss = 0
            assert(len(predict) == len(Y))
            for i in range(len(predict)):
                loss += (predict[i] - Y[i]) ** 2

            #currLoss = loss(X, Y, weights)
            totalLoss.append(loss)

        return np.asarray(weights), np.asarray(totalLoss)

```

Next, we need to load the dataset. In doing so, the following function may be helpful:

```
In [5]: def load_data(filename):
        """
        Function loads data stored in the file filename and returns it as a numpy ndarray.

        Inputs:
            filename: GeneratorExitiven as a string.

        Outputs:
            Data contained in the file, returned as a numpy ndarray
        """
        return np.loadtxt(filename, skiprows=1, delimiter=',')

def getData(data):
    """
    This function takes the raw data and returns two arrays for inputs(x)
    and outputs(y)
    """
    x = []
    y = []
    arr = np.asarray([1.0])
    for i in data:
        x.append(np.concatenate((arr, i[:-1]), axis = 0))
        y.append(i[(len(i) - 1)])
    return np.asarray(x), np.asarray(y)
```

Now, load the dataset in `sgd_data.csv` and run SGD using the given parameters; print out the final weights.

```
In [6]: #=====
        # TODO:
        # (1) load the dataset
        # (2) run SGD using the given parameters
        # (3) print out the final weights.
        #=====
        weights = []
        loss = []
        numEpochs = 800
        epochs = [800, 900, 1000, 1100, 1200, 1300]
        steps = [math.exp(-10), math.exp(-11), math.exp(-12), math.exp(-13), math.exp(-14), math.exp(-15)]
        stp = math.exp(-15)
        allData = load_data("data/sgd_data.csv")
        x, y = getData(allData)

        initial = [0.001, 0.001, 0.001, 0.001, 0.001]
        # X, Y, w_start, eta, N_epochs
        fw, _ = SGD(x, y, initial, stp, numEpochs)
        print(fw)

        for e in epochs:
            for step in steps:
                initial = [0.001, 0.001, 0.001, 0.001, 0.001]
                fw, totalLoss = SGD(x, y, initial, step, numEpochs)
                weights.append(fw)
                loss.append(totalLoss)

        [ -0.22720591  -5.94229011   3.94369494 -11.72402388   8.78549375]
```

Problem 4G: Convergence of SGD


This problem examines the convergence of SGD for different learning rates. Please implement your code in the cell below:

```
In [10]: #=====
# TODO: create a plot showing the convergence
# of SGD for the different learning rates.
#=====
print(loss[0])

fig = plt.figure()
x = epochs
plt.title('Training Error vs. Epochs', fontsize = 22)
plt.plot(loss[0], marker = '.')
plt.plot(loss[1], marker = '.')
plt.plot(loss[2], marker = '.')
plt.plot(loss[3], marker = '.')
plt.plot(loss[4], marker = '.')
plt.plot(loss[5], marker = '.')
plt.legend(('e$^{-10}$ rate', 'e$^{-11}$ rate', 'e$^{-12}$ rate', 'e$^{-13}$ rate',
            'e$^{-14}$ rate', 'e$^{-15}$ rate'), loc = 'upper right', fontsize = 14)
plt.xlabel('Number of Epochs (log scale)', fontsize = 18)
plt.xscale('log')
plt.ylabel('Training Error', fontsize = 18)
```

[2874307.79361655	500820.74432231	116012.3436524	29495.32628487
9933.17467149	5483.21974217	4463.39592301	4227.13315194
4171.40613985	4157.83834506	4154.34642156	4153.36282484
4153.04781185	4152.93004034	4152.87821425	4152.85132614
4152.83483885	4152.82299375	4152.81332182	4152.80471177
4152.79664943	4152.78889313	4152.78132811	4152.77389936
4152.76658012	4152.75935713	4152.75222348	4152.74517535
4152.73821038	4152.73132692	4152.72452366	4152.71779947
4152.71115329	4152.70458412	4152.698091	4152.69167298
4152.68532912	4152.6790585	4152.67286023	4152.6667334
4152.66067713	4152.65469055	4152.64877281	4152.64292305
4152.63714043	4152.63142413	4152.62577334	4152.62018724
4152.61466504	4152.60920596	4152.60380922	4152.59847405
4152.59319971	4152.58798544	4152.58283051	4152.57773419
4152.57269577	4152.56771453	4152.56278978	4152.55792082
4152.55310699	4152.54834759	4152.54364198	4152.53898949
4152.53438948	4152.5298413	4152.52534434	4152.52089795
4152.51650154	4152.51215449	4152.50785621	4152.5036061
4152.49940357	4152.49524806	4152.491139	4152.48707581
4152.48305795	4152.47908487	4152.47515603	4152.47127089
4152.46742893	4152.46362962	4152.45987246	4152.45615693
4152.45248254	4152.44884879	4152.4452552	4152.44170127
4152.43818653	4152.43471053	4152.43127278	4152.42787283
4152.42451024	4152.42118454	4152.41789531	4152.41464211
4152.41142451	4152.40824207	4152.4050944	4152.40198106
4152.39890165	4152.39585578	4152.39284303	4152.38986302
4152.38691535	4152.38399965	4152.38111554	4152.37826264
4152.37544058	4152.372649	4152.36988754	4152.36715584
4152.36445355	4152.36178033	4152.35913583	4152.35651971
4152.35393165	4152.3513713	4152.34883835	4152.34633248
4152.34385336	4152.34140069	4152.33897415	4152.33657343
4152.33419825	4152.33184829	4152.32952326	4152.32722288
4152.32494686	4152.3226949	4152.32046674	4152.31826209
4152.31608069	4152.31392226	4152.31178653	4152.30967325
4152.30758215	4152.30551298	4152.30346548	4152.30143941
4152.29943451	4152.29745054	4152.29548727	4152.29354444
4152.29162184	4152.28971922	4152.28783635	4152.28597301
4152.28412898	4152.28230403	4152.28049795	4152.27871052
4152.27694152	4152.27519075	4152.273458	4152.27174307
4152.27004574	4152.26836582	4152.26670312	4152.26505743
4152.26342857	4152.26181634	4152.26022055	4152.25864103
4152.25707758	4152.25553002	4152.25399818	4152.25248187
4152.25098093	4152.24949517	4152.24802444	4152.24656855
4152.24512735	4152.24370067	4152.24228834	4152.24089021
4152.23950611	4152.2381359	4152.2367794	4152.23543649
4152.23410699	4152.23279076	4152.23148766	4152.23019754
4152.22892025	4152.22765565	4152.2264036	4152.22516397
4152.22393661	4152.22272139	4152.22151818	4152.22032684
4152.21914724	4152.21797926	4152.21682277	4152.21567763
4152.21454373	4152.21342094	4152.21230914	4152.21120821
4152.21011804	4152.20903849	4152.20796947	4152.20691085
4152.20586252	4152.20482437	4152.20379629	4152.20277817
4152.2017699	4152.20077137	4152.19978248	4152.19880313
4152.19783321	4152.19687262	4152.19592126	4152.19497903
4152.19404583	4152.19312157	4152.19220615	4152.19129947
4152.19040144	4152.18951197	4152.18863097	4152.18775835
4152.18689402	4152.18603788	4152.18518986	4152.18434986
4152.1835178	4152.18269361	4152.18187718	4152.18106845
4152.18026732	4152.17947373	4152.17868758	4152.17790881
4152.17713733	4152.17637307	4152.17561595	4152.1748659
4152.17412284	4152.1733867	4152.1726574	4152.17193488
4152.17121907	4152.17050989	4152.16980728	4152.16911117
4152.16842149	4152.16773817	4152.16706115	4152.16639036
4152.16572574	4152.16506722	4152.16441474	4152.16376825
4152.16312767	4152.16249294	4152.16186402	4152.16124083

Out[10]: Text(0, 0.5, 'Training Error')



Problem 4H

Provide your code for computing the least-squares analytical solution below.

```
In [9]: #=====
# TODO: implement the least-squares
# analytical solution.
#=====

x, y = getData(allData)
x = np.matrix(x)
y = np.matrix(y)
ret = np.linalg.inv(x.transpose() * x) * x.transpose() * y.transpose()

print(list(ret))

[matrix([[ -0.31644251]]), matrix([[ -5.99157048]]), matrix([[ 4.01509955]]), matri
x([[ -11.93325972]]), matrix([[ 8.99061096]])]
```

In []: