

1 Basics

Question A: The hypothesis set is a set of hypotheses that attempt to approximate the target function.

Question B: A linear model has the following form for its hypothesis set: $w^T x + b$.

Question C: Overfitting is when test error is much greater than training error; this means that the model trained itself to specifically fit only to the training set (and the set's specific kind of noise) and, when deployed to work on other tests, becomes less accurate.

Question D: One, we can use some kind of k-fold validation to determine which models perform less accurately than others. Two, we can train on more data to reduce levels of variance.

Question E: Training data is what we use to learn and what we use to find parameters of our final model. Test data is what we will use to test our final model to see how accurately it can perform. We should never change our model based on information from test data because doing so might introduce bias into our model that's attuned to the test set.

Question F: We assume that our sampled data has a variety of outputs based on inputs; we also assume that the data is sampled without regards to the actual distribution.

Question G: If want to decide whether or not an email is spam with machine learning, the input space, X can be all of our emails in an organized vector form of some sort. Y could be whether or not the email is spam, 0 or 1 in a binary choice.

Question H: The k-fold cross-validation procedure enables users to reuse data as both training and testing data. K-fold cross-validation occurs when the data set is split into k number of partitions; and k-1 partitions are used as training, with the last partition used as testing. This procedure rotates so that every partition is at sometime used as testing.

2 Bias-Variance Tradeoff

Question A:

Given four definitions,

$$\begin{aligned} F(x) &= E_s[f_s(x)] \\ E_{out}(f_s) &= E_x[(f_s(x) - y(x))^2] \\ Bias(x) &= (F(x) - y(x))^2 \\ Var(x) &= E_s[f_s(x) - F(x)]^2 \end{aligned}$$

we derive the bias-variance decomposition for the squared error loss function in the following way:


We start with

$$E_s[E_{out}(f_s)]$$

and substitute our second definition. Since we see that both the bias and variance function use $F(x)$, we continue with manipulation:

$$\begin{aligned} &= E_s[E_x[(f_s(x) - y(x))^2]] \\ &= E_s[E_x[(f_s(x) - y(x) + F(x) - F(x))^2]] \\ &= E_s[E_x[(f_s(x) - y(x) + F(x) - F(x))^2]] \\ &= E_s[E_x[(f_s(x) - F(x))^2 + (F(x) - y(x))^2 + 2(f_s(x) - F(x))(F(x) - y(x))]] \\ &= E_x[E_s[(f_s(x) - F(x))^2 + (F(x) - y(x))^2 + 2(f_s(x) - F(x))(F(x) - y(x))]] \\ &= E_x[E_s[(f_s(x) - F(x))^2]] + E_x[(F(x) - y(x))^2] \\ &= [Bias(x) + Var(x)] \end{aligned}$$

Question B:



Question C:

Since the graph for D1 (degree 1) has a higher average error compared to the other graphs, it must have the bias. Also, we know from lecture that bias decreases with model complexity so it makes sense that the low complexity of the model means it has high bias and underfits the model.

Question D:

The model with highest variance should be the D12 (degree 12) model because it has low training error but large validation error at earlier data points. This makes sense since this model has a high complexity, which would correlate with overfitting and high variance.

Question E:

The learning curve of the quadratic model tells us the model would predict more accurately with more data; we see that the graph stabilizes out after the number of data points increases.

Question F:

Training error is generally lower than validation error because our model is trained on the training data and therefore is more attuned to the training data. We never train on our validation error so it makes sense so that there would be higher error on our validation data, making training error lower than validation error.


Question G:

Based on the learning curves, I would expect D6 (degree 6) model to perform best on unseen data since it seems to have lower validation errors in regards to larger amounts of data points.

3 The Perceptron

Question A:

t	b	w1	w2
0	0.0	0.0	1.0
1	1.0	1.0	-1.0
2	2.0	1.0	2.0
3	3.0	2.0	0.0
4	3.0	2.0	0.0



Question B: In a 2D dataset, the smallest dataset that is not linearly separable has 4 data points that resembles

$\begin{array}{c} + \\ - \quad - \end{array}$
 $\begin{array}{c} + \\ - \quad - \end{array}$

A 3D dataset would have a similar arrangement of points, except in 3D space and with 5 points. More generally, this means that for an N-dimensional case, the smallest data set that

is NOT linearly separable has $N+2$ points.

Question C: If a dataset is not linearly separable, then there exists no way for a linear model to correctly identify/classify all the data. Therefore, based on the PL algorithm, there will never be an iteration where all datapoints are correctly classified so the algorithm will continue running and never converge.

4 Stochastic Gradient Descent

Question A:

To define x and w such that the model includes the bias term, we can include a bias term as the zeroth values, as $w_0 = b$ and $x_0 = 1$. So we have $w = (w_0, w_1, \dots, w_d)$ and $x = (1, x_1, x_2, \dots, x_d)$

Question B:


$$\sum_{i=1}^N -2(y_i - w^T x_i) x_i$$

Question C:

A sample of weights.


```

0: 43.08446558593781, [ 0.01161911 0.00960535]
1: 42.80858797312999, [ 0.01323284 0.00921207]
2: 42.534545680530826, [ 0.01484121 0.00882016]
3: 42.26232649682018, [ 0.01644424 0.00842962]
4: 41.991918291935725, [ 0.01804194 0.00804044]
5: 41.72330901653238, [ 0.01963433 0.00765262]
6: 41.456486701444874, [ 0.02122143 0.00726615]
7: 41.191439457154495, [ 0.02280326 0.00688102]
8: 40.928155473258386, [ 0.02437993 0.00649724]
9: 40.66623017943884, [ 0.02595116 0.0061148 ]
10: 40.4068304374648, [ 0.02751727 0.00573369]
11: 40.14876615562224, [ 0.02907818 0.00535391]
12: 39.8924186732485, [ 0.0306339 0.00497546]
13: 39.6377765769436, [ 0.03218445 0.00459832]
14: 39.384828492319606, [ 0.03372985 0.0042225 ]
15: 39.133563175987646, [ 0.03527012 0.003848 ]
16: 38.88396942256283, [ 0.03680527 0.00347479]
17: 38.63603611041127, [ 0.03833531 0.00310289]
18: 38.38975219190503, [ 0.03986027 0.00273228]
19: 38.145106692929716, [ 0.04139017 0.00236297]
20: 37.90208871239513, [ 0.04295052 0.00199495]
21: 37.6606874217931, [ 0.04440483 0.00162821]
22: 37.42089206449599, [ 0.04590963 0.00126274]
23: 37.18269195571486, [ 0.04740942 0.00089856]
24: 36.94607648158497, [ 0.04890432 0.00053564]
25: 36.711035098912426, [ 0.05039408 0.00017399]
26: 36.47755733465941, [ 0.05187998 0.0001864 ]
27: 36.24563278547776, [ 0.05335894 0.00054554]
28: 36.0152511172453, [ 0.05483398 0.00090342]
29: 35.7864020646049, [ 0.05630413 0.00126005]
30: 35.55907543050686, [ 0.05776939 0.00161544]
31: 35.3326108575444, [ 0.05922978 0.00196959]
32: 35.10894896855218, [ 0.06068521 0.0023225 ]
33: 34.88612908405768, [ 0.06213602 0.00267418]
34: 34.66479150393524, [ 0.06358104 0.00302464]
35: 34.44492636591486, [ 0.06502298 0.00337387]
36: 34.226523873350246, [ 0.06645927 0.00372189]
37: 34.00957429478419, [ 0.06789079 0.00406869]
38: 33.79406796351348, [ 0.06931755 0.00441428]
39: 33.579995277158865, [ 0.07073957 0.00475867]
40: 33.36734669723621, [ 0.07215686 0.00510185]
41: 33.15611274873163, [ 0.07356945 0.00544384]
42: 32.94628401967916, [ 0.07497734 0.00578463]
43: 32.73785116074099, [ 0.07638056 0.00612424]
44: 32.53080488479094, [ 0.07777911 0.00646266]
45: 32.32513596650003, [ 0.07917302 0.0067989 ]
46: 32.120835241925846, [ 0.0805623 0.00713596]
47: 31.91789360810344, [ 0.08194695 0.00747085]
48: 31.71630202639127, [ 0.08332701 0.00780518]
49: 31.516051503310642, [ 0.08470249 0.00813714]
50: 31.317133127660234, [ 0.08607339 0.00846853]
51: 31.11953803260334, [ 0.08743974 0.00879877]
52: 30.92325741403041, [ 0.08880155 0.00912786]
53: 30.72828256414773, [ 0.09015883 0.0094558 ]
54: 30.534604682423076, [ 0.09151161 0.0097826 ]
55: 30.34221525257957, [ 0.09285989 0.01010826]
56: 30.15110566462312, [ 0.0942037 0.01043278]
57: 30.00000000000000, [ 0.00000000 0.00000000]
58: 30.00000000000000, [ 0.00000000 0.00000000]
59: 30.00000000000000, [ 0.00000000 0.00000000]
60: 30.00000000000000, [ 0.00000000 0.00000000]
61: 30.00000000000000, [ 0.00000000 0.00000000]
62: 30.00000000000000, [ 0.00000000 0.00000000]
63: 30.00000000000000, [ 0.00000000 0.00000000]
64: 30.00000000000000, [ 0.00000000 0.00000000]
65: 30.00000000000000, [ 0.00000000 0.00000000]
66: 30.00000000000000, [ 0.00000000 0.00000000]
67: 30.00000000000000, [ 0.00000000 0.00000000]
68: 30.00000000000000, [ 0.00000000 0.00000000]
69: 30.00000000000000, [ 0.00000000 0.00000000]
70: 30.00000000000000, [ 0.00000000 0.00000000]
71: 30.00000000000000, [ 0.00000000 0.00000000]
72: 30.00000000000000, [ 0.00000000 0.00000000]
73: 30.00000000000000, [ 0.00000000 0.00000000]
74: 30.00000000000000, [ 0.00000000 0.00000000]
75: 30.00000000000000, [ 0.00000000 0.00000000]
76: 30.00000000000000, [ 0.00000000 0.00000000]
77: 30.00000000000000, [ 0.00000000 0.00000000]
78: 30.00000000000000, [ 0.00000000 0.00000000]
79: 30.00000000000000, [ 0.00000000 0.00000000]
80: 30.00000000000000, [ 0.00000000 0.00000000]
81: 30.00000000000000, [ 0.00000000 0.00000000]
82: 30.00000000000000, [ 0.00000000 0.00000000]
83: 30.00000000000000, [ 0.00000000 0.00000000]
84: 30.00000000000000, [ 0.00000000 0.00000000]
85: 30.00000000000000, [ 0.00000000 0.00000000]
86: 30.00000000000000, [ 0.00000000 0.00000000]
87: 30.00000000000000, [ 0.00000000 0.00000000]
88: 30.00000000000000, [ 0.00000000 0.00000000]
89: 30.00000000000000, [ 0.00000000 0.00000000]
90: 30.00000000000000, [ 0.00000000 0.00000000]
91: 30.00000000000000, [ 0.00000000 0.00000000]
92: 30.00000000000000, [ 0.00000000 0.00000000]
93: 30.00000000000000, [ 0.00000000 0.00000000]
94: 30.00000000000000, [ 0.00000000 0.00000000]
95: 30.00000000000000, [ 0.00000000 0.00000000]
96: 30.00000000000000, [ 0.00000000 0.00000000]
97: 30.00000000000000, [ 0.00000000 0.00000000]
98: 30.00000000000000, [ 0.00000000 0.00000000]
99: 30.00000000000000, [ 0.00000000 0.00000000]
100: 30.00000000000000, [ 0.00000000 0.00000000]
101: 30.00000000000000, [ 0.00000000 0.00000000]
102: 30.00000000000000, [ 0.00000000 0.00000000]
103: 30.00000000000000, [ 0.00000000 0.00000000]
104: 30.00000000000000, [ 0.00000000 0.00000000]
105: 30.00000000000000, [ 0.00000000 0.00000000]
106: 30.00000000000000, [ 0.00000000 0.00000000]
107: 30.00000000000000, [ 0.00000000 0.00000000]
108: 30.00000000000000, [ 0.00000000 0.00000000]
109: 30.00000000000000, [ 0.00000000 0.00000000]
110: 30.00000000000000, [ 0.00000000 0.00000000]
111: 30.00000000000000, [ 0.00000000 0.00000000]
112: 30.00000000000000, [ 0.00000000 0.00000000]
113: 30.00000000000000, [ 0.00000000 0.00000000]
114: 30.00000000000000, [ 0.00000000 0.00000000]
115: 30.00000000000000, [ 0.00000000 0.00000000]
116: 30.00000000000000, [ 0.00000000 0.00000000]
117: 30.00000000000000, [ 0.00000000 0.00000000]
118: 30.00000000000000, [ 0.00000000 0.00000000]
119: 30.00000000000000, [ 0.00000000 0.00000000]
120: 30.00000000000000, [ 0.00000000 0.00000000]
121: 30.00000000000000, [ 0.00000000 0.00000000]
122: 30.00000000000000, [ 0.00000000 0.00000000]
123: 30.00000000000000, [ 0.00000000 0.00000000]
124: 30.00000000000000, [ 0.00000000 0.00000000]
125: 30.00000000000000, [ 0.00000000 0.00000000]
126: 30.00000000000000, [ 0.00000000 0.00000000]
127: 30.00000000000000, [ 0.00000000 0.00000000]
128: 30.00000000000000, [ 0.00000000 0.00000000]
129: 30.00000000000000, [ 0.00000000 0.00000000]
130: 30.00000000000000, [ 0.00000000 0.00000000]
131: 30.00000000000000, [ 0.00000000 0.00000000]
132: 30.00000000000000, [ 0.00000000 0.00000000]
133: 30.00000000000000, [ 0.00000000 0.00000000]
134: 30.00000000000000, [ 0.00000000 0.00000000]
135: 30.00000000000000, [ 0.00000000 0.00000000]
136: 30.00000000000000, [ 0.00000000 0.00000000]
137: 30.00000000000000, [ 0.00000000 0.00000000]
138: 30.00000000000000, [ 0.00000000 0.00000000]
139: 30.00000000000000, [ 0.00000000 0.00000000]
140: 30.00000000000000, [ 0.00000000 0.00000000]
141: 30.00000000000000, [ 0.00000000 0.00000000]
142: 30.00000000000000, [ 0.00000000 0.00000000]
143: 30.00000000000000, [ 0.00000000 0.00000000]
144: 30.00000000000000, [ 0.00000000 0.00000000]
145: 30.00000000000000, [ 0.00000000 0.00000000]
146: 30.00000000000000, [ 0.00000000 0.00000000]
147: 30.00000000000000, [ 0.00000000 0.00000000]
148: 30.00000000000000, [ 0.00000000 0.00000000]
149: 30.00000000000000, [ 0.00000000 0.00000000]
150: 30.00000000000000, [ 0.00000000 0.00000000]
151: 30.00000000000000, [ 0.00000000 0.00000000]
152: 30.00000000000000, [ 0.00000000 0.00000000]
153: 30.00000000000000, [ 0.00000000 0.00000000]
154: 30.00000000000000, [ 0.00000000 0.00000000]
155: 30.00000000000000, [ 0.00000000 0.00000000]
156: 30.00000000000000, [ 0.00000000 0.00000000]
157: 30.00000000000000, [ 0.00000000 0.00000000]
158: 30.00000000000000, [ 0.00000000 0.00000000]
159: 30.00000000000000, [ 0.00000000 0.00000000]
160: 30.00000000000000, [ 0.00000000 0.00000000]
161: 30.00000000000000, [ 0.00000000 0.00000000]
162: 30.00000000000000, [ 0.00000000 0.00000000]
163: 30.00000000000000, [ 0.00000000 0.00000000]
164: 30.00000000000000, [ 0.00000000 0.00000000]
165: 30.00000000000000, [ 0.00000000 0.00000000]
166: 30.00000000000000, [ 0.00000000 0.00000000]
167: 30.00000000000000, [ 0.00000000 0.00000000]
168: 30.00000000000000, [ 0.00000000 0.00000000]
169: 30.00000000000000, [ 0.00000000 0.00000000]
170: 30.00000000000000, [ 0.00000000 0.00000000]
171: 30.00000000000000, [ 0.00000000 0.00000000]
172: 30.00000000000000, [ 0.00000000 0.00000000]
173: 30.00000000000000, [ 0.00000000 0.00000000]
174: 30.00000000000000, [ 0.00000000 0.00000000]
175: 30.00000000000000, [ 0.00000000 0.00000000]
176: 30.00000000000000, [ 0.00000000 0.00000000]
177: 30.00000000000000, [ 0.00000000 0.00000000]
178: 30.00000000000000, [ 0.00000000 0.00000000]
179: 30.00000000000000, [ 0.00000000 0.00000000]
180: 30.00000000000000, [ 0.00000000 0.00000000]
181: 30.00000000000000, [ 0.00000000 0.00000000]
182: 30.00000000000000, [ 0.00000000 0.00000000]
183: 30.00000000000000, [ 0.00000000 0.00000000]
184: 30.00000000000000, [ 0.00000000 0.00000000]
185: 30.00000000000000, [ 0.00000000 0.00000000]
186: 30.00000000000000, [ 0.00000000 0.00000000]
187: 30.00000000000000, [ 0.00000000 0.00000000]
188: 30.00000000000000, [ 0.00000000 0.00000000]
189: 30.00000000000000, [ 0.00000000 0.00000000]
190: 30.00000000000000, [ 0.00000000 0.00000000]
191: 30.00000000000000, [ 0.00000000 0.00000000]
192: 30.00000000000000, [ 0.00000000 0.00000000]
193: 30.00000000000000, [ 0.00000000 0.00000000]
194: 30.00000000000000, [ 0.00000000 0.00000000]
195: 30.00000000000000, [ 0.00000000 0.00000000]
196: 30.00000000000000, [ 0.00000000 0.00000000]
197: 30.00000000000000, [ 0.00000000 0.00000000]
198: 30.00000000000000, [ 0.00000000 0.00000000]
199: 30.00000000000000, [ 0.00000000 0.00000000]
200: 30.00000000000000, [ 0.00000000 0.00000000]
201: 30.00000000000000, [ 0.00000000 0.00000000]
202: 30.00000000000000, [ 0.00000000 0.00000000]
203: 30.00000000000000, [ 0.00000000 0.00000000]
204: 30.00000000000000, [ 0.00000000 0.00000000]
205: 30.00000000000000, [ 0.00000000 0.00000000]
206: 30.00000000000000, [ 0.00000000 0.00000000]
207: 30.00000000000000, [ 0.00000000 0.00000000]
208: 30.00000000000000, [ 0.00000000 0.00000000]
209: 30.00000000000000, [ 0.00000000 0.00000000]
210: 30.00000000000000, [ 0.00000000 0.00000000]
211: 30.00000000000000, [ 0.00000000 0.00000000]
212: 30.00000000000000, [ 0.00000000 0.00000000]
213: 30.00000000000000, [ 0.00000000 0.00000000]
214: 30.00000000000000, [ 0.00000000 0.00000000]
215: 30.00000000000000, [ 0.00000000 0.00000000]
216: 30.00000000000000, [ 0.00000000 0.00000000]
217: 30.00000000000000, [ 0.00000000 0.00000000]
218: 30.00000000000000, [ 0.00000000 0.00000000]
219: 30.00000000000000, [ 0.00000000 0.00000000]
220: 30.00000000000000, [ 0.00000000 0.00000000]
221: 30.00000000000000, [ 0.00000000 0.00000000]
222: 30.00000000000000, [ 0.00000000 0.00000000]
223: 30.00000000000000, [ 0.00000000 0.00000000]
224: 30.00000000000000, [ 0.00000000 0.00000000]
225: 30.00000000000000, [ 0.00000000 0.00000000]
226: 30.00000000000000, [ 0.00000000 0.00000000]
227: 30.00000000000000, [ 0.00000000 0.00000000]
228: 30.00000000000000, [ 0.00000000 0.00000000]
229: 30.00000000000000, [ 0.00000000 0.00000000]
230: 30.00000000000000, [ 0.00000000 0.00000000]
231: 30.00000000000000, [ 0.00000000 0.00000000]
232: 30.00000000000000, [ 0.00000000 0.00000000]
233: 30.00000000000000, [ 0.00000000 0.00000000]
234: 30.00000000000000, [ 0.00000000 0.00000000]
235: 30.00000000000000, [ 0.00000000 0.00000000]
236: 30.00000000000000, [ 0.00000000 0.00000000]
237: 30.00000000000000, [ 0.00000000 0.00000000]
238: 30.00000000000000, [ 0.00000000 0.00000000]
239: 30.00000000000000, [ 0.00000000 0.00000000]
240: 30.00000000000000, [ 0.00000000 0.00000000]
241: 30.00000000000000, [ 0.00000000 0.00000000]
242: 30.00000000000000, [ 0.00000000 0.00000000]
243: 30.00000000000000, [ 0.00000000 0.00000000]
244: 30.00000000000000, [ 0.00000000 0.00000000]
245: 30.00000000000000, [ 0.00000000 0.00000000]
246: 30.00000000000000, [ 0.00000000 0.00000000]
247: 30.00000000000000, [ 0.00000000 0.00000000]
248: 30.00000000000000
```

**Question F:**

After running my code, my final weights were:

`[-0.22720591 -5.94229011 3.94369494 -11.72402388 8.78549375]`

Question G:**Question H:**

Computing this analytical solution gives me: `matrix([-0.31644251])`, `matrix([-5.99157048])`, `matrix([4.01509955])`, `matrix([-11.93325972])`, `matrix([8.99061096])`.

Which seems close to our previously generated solution of `[-0.22720591 -5.94229011 3.94369494 -11.72402388 8.78549375]`

Question I:

We would use SGD when a closed form solution exists if we are given some very large input data and using closed form solution would take time and resources. SGD is faster to compute and would be the preferred method since it's less computationally complex as data grows in

size.

Question J:

Based on the SGD convergence plots generated earlier, a stopping condition more sophisticated than a pre-defined number of epochs would be finding convergence points for certain η s and stopping around those pre-defined convergence points.

Question K:

The convergence behavior of the weight vector differs between the perceptron and SGD algorithms in that SGD convergence behavior varies significantly based on the size of steps used in the algorithm.

Problem 2

```
In [3]: # Setup:

import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import KFold

import warnings
warnings.filterwarnings("ignore")
```

Example code using the polyfit and kfold functions

Note: This section is not part of the homework problem, but provides some potentially-helpful example code regarding the usage of `numpy.polyfit`, `numpy.polyval`, and `sklearn.model_selection.KFold`.

First, let's generate some synthetic data: a quadratic function plus some Gaussian noise.


```
In [4]: # Coefficients of the quadratic function,  $y(x) = ax^2 + bx + c$ :
a = 2
b = 5
c = 7

N = 200          # Number of data points
x = np.linspace(-10, 10, num = N)          # x ranges from -10 to 10
# y is the quadratic function of x specified by a, b, and c, plus noise
y = a*x**2 + b*x + c + 15* np.random.randn(N)

# Plot the data:
plt.figure()
plt.plot(x, y, marker = 'o', linewidth = 0)

plt.xlabel('x')
plt.ylabel('y')
plt.title('Synthetic data')

plt.show()
```



Next, we'll use the `numpy.polyfit` function to fit a quadratic polynomial to this data. We can evaluate the resulting polynomial at arbitrary points.

```
In [5]: # Fit a degree-2 polynomial to the data:
degree = 2
coefficients = np.polyfit(x, y, degree)

# Print out the resulting quadratic function:
print('We fit the following quadratic function: f(x) = %f*x^2 + %f*x + %f' % \
      (coefficients[0], coefficients[1], coefficients[2]))

# Evaluate the fitted polynomial at x = 4:
x_test = 4
f_eval = np.polyval(coefficients, x_test)
print('\nf(%i) = %f' % (x_test, f_eval))

# Let's visualize our fitted quadratic:
plt.figure()


plt.plot(x, y, marker = 'o', linewidth = 0)
plt.plot(x, np.polyval(coefficients, x), color = 'red', linewidth = 3)

plt.legend(['Data', 'Fitted quadratic'], loc = 'best')
plt.xlabel('x')
plt.ylabel('y')
plt.title('Synthetic data with Fitted Quadratic')

plt.show()
```

We fit the following quadratic function: $f(x) = 1.985406x^2 + 5.113636x + 7.600382$

$f(4) = 59.821414$



Finally, assume that we'd like to perform 10-fold cross validation with this dataset. Let's divide it into training and test sets, and print out the test sets. To limit the amount of text that we are printing out, we'll modify the dataset to make it smaller.

```
In [6]: # Coefficients of the quadratic function,  $y = ax^2 + bx + c$ :
a = 2
b = 5
c = 7

N = 80          # Number of points--fewer this time!
x = np.linspace(-10, 10, num = N)          # x ranges from -10 to 10
# y is the quadratic function of x specified by a, b, and c, plus noise
y = a*x**2 + b*x + c + 15* np.random.randn(N)

# Initialize kfold cross-validation object with 10 folds:
num_folds = 10
kf = KFold(n_splits=num_folds)

# Iterate through cross-validation folds:
i = 1
for train_index, test_index in kf.split(x):

    # Print out test indices:
    print('Fold ', i, ' of ', num_folds, ' test indices:', test_index)

    # Training and testing data points for this fold:
    x_train, x_test = x[train_index], x[test_index]
    y_train, y_test = y[train_index], y[test_index]

    i += 1
```

```
Fold 1 of 10 test indices: [0 1 2 3 4 5 6 7]
Fold 2 of 10 test indices: [ 8  9 10 11 12 13 14 15]
Fold 3 of 10 test indices: [16 17 18 19 20 21 22 23]
Fold 4 of 10 test indices: [24 25 26 27 28 29 30 31]
Fold 5 of 10 test indices: [32 33 34 35 36 37 38 39]
Fold 6 of 10 test indices: [40 41 42 43 44 45 46 47]
Fold 7 of 10 test indices: [48 49 50 51 52 53 54 55]
Fold 8 of 10 test indices: [56 57 58 59 60 61 62 63]
Fold 9 of 10 test indices: [64 65 66 67 68 69 70 71]
Fold 10 of 10 test indices: [72 73 74 75 76 77 78 79]
```

Loading the Data for Problem 2

This code loads the data from `bv_data.csv` using the `load_data` helper function. Note that `data[:, 0]` is an array of all the `x` values in the data and `data[:, 1]` is an array of the corresponding `y` values.

```
In [7]: def load_data(filename):
        """
        Function loads data stored in the file filename and returns it as a numpy ndarray.
        Input:
            filename: given as a string.
        Output:
            Data contained in the file, returned as a numpy ndarray
        """
        return np.loadtxt(filename, skiprows=1, delimiter=',')
```

```
In [8]: data = load_data('data/bv_data.csv')
x = data[:, 0]
y = data[:, 1]
```

Write your code below for solving problem 2 part B:

```

In [9]: def sqErr(n1, n2):
        # return square error
        return (n1 - n2) ** 2

def getData(data):
    # return arrays representing x and y
    x = []
    y = []

    for i in data:
        x.append(i[0])
        y.append(i[1])

    return np.asarray(x), np.asarray(y)

def retKFoldData(n, numFolds, data):
    # return training data and validation data for one fold
    # fold number = section number that serves as validation number

    foldNumber = int(n)
    part = int(len(data) / numFolds)
    numFolds = int(numFolds)

    lower = (foldNumber - 1) * part
    upper = foldNumber * part

    training = np.concatenate((data[:lower], data[upper:]), axis = 0)
    validation = data[lower : upper]
    return training, validation

def avgErr(data, folds, degreePoly):
    # return avg training & validation errors

    trainingErrs = []
    validationErrs = []

    for k in range(1, folds + 1):
        trainingErr = 0
        validationErr = 0
        training, validation = retKFoldData(k, folds, data) #Get training and validation data from the retKFoldData function

        trainX, trainY = getData(training)
        validationX, validationY = getData(validation)
        z = np.polyfit(trainX, trainY, degreePoly) #np.polyfit will return coefficients

        # use training data to find training error
        for a in range(len(trainX)):
            p1 = np.polyval(z, trainX[a])
            trainingErr += sqErr(p1, trainY[a])

        # use validation data to find validation error
        for b in range(len(validationX)):
            p2 = np.polyval(z, validationX[b])
            validationErr += sqErr(p2, validationY[b])

        trainingErrs.append(trainingErr / len(trainX))
        validationErrs.append(validationErr / len(validationX))

    return sum(trainingErrs) / folds, sum(validationErrs) / folds

```

```

In [10]: finalData = load_data("data/bv_data.csv")

trainingErr = []
validationErr = []

for d in [1, 2, 6, 12]:
    for size in range(20,101,5):
        t = finalData[:size]
        avgTrain, avgValidation = avgErr(t, 5, d)

        trainingErr.append(avgTrain)
        validationErr.append(avgValidation)

train1 = trainingErr[:17]
valid1 = validationErr[:17]

train2 = trainingErr[17:34]
valid2 = validationErr[17:34]

train6 = trainingErr[34: 51]
valid6 = validationErr[34:51]

train12 = trainingErr[51:]
valid12 = validationErr[51:]

# plotting figures
N = list(range(20,101,5))
x = N

plt.figure(1)
plt.subplot(221)

plt.title('D1')
plt.legend(('Training', 'Validation'))
plt.plot(x, train1, x, valid1, marker = '.')
plt.ylabel('Avg Error')
plt.margins(tight=True)

plt.subplot(222)
plt.title('D2')
plt.plot(x, train2, x, valid2, marker = '.')
plt.margins(tight=True)


plt.subplot(223)
plt.title('D6')
plt.plot(x, train6, x, valid6, marker = '.')
plt.xlabel('Num Data')
plt.ylabel('Avg Error')
plt.margins(tight=True)

plt.subplot(224)
plt.title('D12')
plt.plot(x, train12, x, valid12, marker = '.')
plt.xlabel('Num Data')
# use log scale for fit
plt.yscale('log')
plt.margins(tight=True)

# plt.subplots_adjust(left=.5, bottom=.5, right=.5, top=.5, wspace=.5, hspace=None)

```

Out[10]: (0.05, 0.05)



In []:

In []:

In []:

In []:

In []:

Problem 3

Use this notebook to write your code for problem 3 by filling in the sections marked `# TODO` and running all cells.

```
In [61]: import numpy as np
import matplotlib.pyplot as plt
import itertools
from prettytable import PrettyTable

from perceptron_helper import (
    predict,
    plot_data,
    boundary,
    plot_perceptron,
)

%matplotlib inline
```

Implementation of Perceptron

First, we will implement the perceptron algorithm. Fill in the `update_perceptron()` function so that it finds a single misclassified point and updates the weights and bias accordingly. If no point exists, the weights and bias should not change.

Hint: You can use the `predict()` helper method, which labels a point 1 or -1 depending on the weights and bias.


```
In [62]: def update_perceptron(X, Y, w, b):
        """
        This method updates a perceptron model. Takes in the previous weights
        and returns weights after an update, which could be nothing.

        Inputs:
            X: A (N, D) shaped numpy array containing N points of D dimensions.
            Y: A (N, ) shaped numpy array containing the N labels (one for each point
            in X).
            w: A (D, ) shaped numpy array containing the initial weight vector of D di
            mensions.
            b: A float containing the bias term.

        Output:
            next_w: A (D, ) shaped numpy array containing the next weight vector
            after updating on a single misclassified point, if one exists.
            next_b: The next float bias term after updating on a single
            misclassified point, if one exists.
        """
        next_w, next_b = np.copy(w), np.copy(b)

        #=====
        # TODO: Implement update rule for perceptron.
        #=====

        for i in range(len(X)):
            # if misclassified
            if predict(X[i], w, b) != Y[i]:
                for j in range(len(X[i])):
                    # adjust weights for every dimension
                    next_w[j] = w[j] + Y[i]*X[i][j]
                    next_b = b + Y[i]

        return next_w, next_b
```

Next you will fill in the `run_perceptron()` method. The method performs single updates on a misclassified point until convergence, or `max_iter` updates are made. The function will return the final weights and bias. You should use the `update_perceptron()` method you implemented above.

```
In [63]: def run_perceptron(X, Y, w, b, max_iter):
    """
    This method runs the perceptron learning algorithm. Takes in initial weights
    and runs max_iter update iterations. Returns final weights and bias.

    Inputs:
        X: A (N, D) shaped numpy array containing N points of D dimensions.
        Y: A (N, ) shaped numpy array containing the N labels (one for each point
    in X).
        w: A (D, ) shaped numpy array containing the initial weight vector of D di
    mensions.
        b: A float containing the initial bias term.
        max_iter: An int for the maximum number of updates evaluated.

    Output:
        w: A (D, ) shaped numpy array containing the final weight vector.
        b: The final float bias term.
    """

    #=====
    # TODO: Implement perceptron update loop.
    #=====

    for _ in range(max_iter):
        w, b = update_perceptron(X, Y, w, b)

    return w, b
```


Problem 3A

Visualizing a Toy Dataset

We will begin by training our perceptron on a toy dataset of 3 points. The green points are labelled +1 and the red points are labelled -1. We use the helper function `plot_data()` to do so.

```
In [64]: X = np.array([[ -3, -1], [0, 3], [1, -2]])
        Y = np.array([ -1, 1, 1])
```

```
In [65]: fig = plt.figure(figsize=(5,4))
ax = fig.gca(); ax.set_xlim(-4.1, 3.1); ax.set_ylim(-3.1, 4.1)
plot_data(X, Y, ax)
```



Running the Perceptron

Next, we will run the perceptron learning algorithm on this dataset. Update the code to show the weights and bias at each timestep and the misclassified point used in each update. You may change the `update_perceptron()` method to do this, but be sure to update the starter code as well to reflect those changes.

Run the below code, and fill in the corresponding table in the set.

```
In [66]: # Initialize weights and bias.
weights = np.array([0.0, 1.0])
bias = 0.0

weights, bias = run_perceptron(X, Y, weights, bias, 16)

print()
print ("final w = %s, final b = %.1f" % (weights, bias))

final w = [2. 0.], final b = 3.0
```

Visualizing the Perceptron

Getting all that information in table form isn't very informative. Let us visualize what the decision boundaries are at each timestep instead.

The helper functions `boundary()` and `plot_perceptron()` plot a decision boundary given a perceptron weights and bias. Note that the equation for the decision boundary is given by:

$$w_1x_1 + w_2x_2 + b = 0.$$

Using some algebra, we can obtain x_2 from x_1 to plot the boundary as a line.

$$x_2 = \frac{-w_1x_1 - b}{w_2}.$$

Below is a redefinition of the `run_perceptron()` method to visualize the points and decision boundaries at each timestep instead of printing. Fill in the method using your previous `run_perceptron()` method, and the above helper methods.

Hint: The `axs` element is a list of Axes, which are used as subplots for each timestep. You can do the following:

```
ax = axs[i]
```

to get the plot corresponding to $t = i$. You can then use `ax.set_title()` to title each subplot. You will want to use the `plot_data()` and `plot_perceptron()` helper methods.

```
In [78]: def run_perceptron(X, Y, w, b, axs, max_iter):
    """
    This method runs the perceptron learning algorithm. Takes in initial weights
    and runs max_iter update iterations. Returns final weights and bias.

    Inputs:
        X: A (N, D) shaped numpy array containing N points of D dimensions.
        Y: A (N, ) shaped numpy array containing the N labels (one for each point
    in X).
        w: A (D, ) shaped numpy array containing the initial weight vector of D di
    mensions.
        b: A float containing the initial bias term.
        axs: A list of Axes that contain suplots for each timestep.
        max_iter: An int for the maximum number of updates evaluated.

    Output:
        The final weight and bias vectors.
    """

    #=====
    # TODO: Implement perceptron update loop.
    #=====
    time = 0
    t = PrettyTable()

    t.field_names = ["t", "b", "w1", "w2"]
    t.add_row([time, b, w[0], w[1]])

    for i in range(max_iter):
        time += 1
        w, b = update_perceptron(X, Y, w, b)
        ax = axs[i]
        title = "Iteration " + str(i)
        ax.set_title(title)
        plot_perceptron(w, b, ax)
        plot_data(X, Y, ax)
        t.add_row([time, b, w[0], w[1]])

    print(t)
    return w, b
```

Run the below code to get a visualization of the perceptron algorithm. The red region are areas the perceptron thinks are negative examples.


```
In [79]: # Initialize weights and bias.
weights = np.array([0.0, 1.0])
bias = 0.0

# Note: there are 4 subplots and max_iter=4 below. Plot BEFORE each timestep update.
f, ax_arr = plt.subplots(2, 2, sharex=True, sharey=True, figsize=(9,8))
axs = list(itertools.chain.from_iterable(ax_arr))
for ax in axs:
    ax.set_xlim(-4.1, 3.1); ax.set_ylim(-3.1, 4.1)

run_perceptron(X, Y, weights, bias, axs, 4)

f.tight_layout()
```

t	b	w1	w2
0	0.0	0.0	1.0
1	1.0	1.0	-1.0
2	2.0	1.0	2.0
3	3.0	2.0	0.0
4	3.0	2.0	0.0




Problem 3C

Visualize a Non-linearly Separable Dataset.

We will now work on a dataset that cannot be linearly separated, namely one that is generated by the XOR function.

```
In [33]: X = np.array([[0, 1], [1, 0], [0, 0], [1, 1]])  
Y = np.array([1, 1, -1, -1])
```

```
In [34]: fig = plt.figure(figsize=(5,4))  
ax = fig.gca(); ax.set_xlim(-0.1, 1.1); ax.set_ylim(-0.1, 1.1)  
plot_data(X, Y, ax)
```



We will now run the perceptron algorithm on this dataset. We will limit the total timesteps this time, but you should see a pattern in the updates. Run the below code.

```


In [35]: # Initialize weights and bias.
weights = np.array([0.0, 1.0])
bias = 0.0

# Note: there are 16 subplots and max_iter=16. Plot BEFORE each timestep update.
f, ax_arr = plt.subplots(4, 4, sharex=True, sharey=True, figsize=(9,8))
axs = list(itertools.chain.from_iterable(ax_arr))
for ax in axs:
    ax.set_xlim(-0.1, 1.1); ax.set_ylim(-0.1, 1.1)

run_perceptron(X, Y, weights, bias, axs, 16)

f.tight_layout()

```



In []:

Problem 4, Parts C-E: Stochastic Gradient Descent Visualization

In this Jupyter notebook, we visualize how SGD works. This visualization corresponds to parts C-E of question 4 in set 1.

Use this notebook to write your code for problem 4 parts C-E by filling in the sections marked `# TODO` and running all cells.

```
In [81]: # Setup.

import numpy as np
import matplotlib.pyplot as plt
from IPython.display import HTML
import math

from sgd_helper import (
    generate_dataset1,
    generate_dataset2,
    plot_dataset,
    plot_loss_function,
    animate_convergence,
    animate_sgd_suite
)
```

Problem 4C: Implementation of SGD

Fill in the loss, gradient, and SGD functions according to the guidelines given in the problem statement in order to perform SGD.


```

In [82]: def loss(X, Y, w):
    """
    Calculate the squared loss function.

    Inputs:
        X: A (N, D) shaped numpy array containing the data points.
        Y: A (N, ) shaped numpy array containing the (float) labels of the data points.
        w: A (D, ) shaped numpy array containing the weight vector.

    Outputs:
        The loss evaluated with respect to X, Y, and w.
    """
    predict = []
    for x in X:
        predict.append(np.inner(w, x))
    predict = np.asarray(predict)

    loss = 0
    for i in range(len(predict)):
        loss += (predict[i] - Y[i]) ** 2

    return loss

def gradient(x, y, w):
    """
    Calculate the gradient of the loss function with respect to
    a single point (x, y), and using weight vector w.

    Inputs:
        x: A (D, ) shaped numpy array containing a single data point.
        y: The float label for the data point.
        w: A (D, ) shaped numpy array containing the weight vector.

    Output:
        The gradient of the loss with respect to x, y, and w.
    """
    #=====
    # TODO: Implement the gradient of the loss function.
    #=====
    grad = -2 * (y - np.inner(w, x)) * x
    return grad

```

```

In [ ]: def SGD(X, Y, w_start, eta, N_epochs):
    """
    Perform SGD using dataset (X, Y), initial weight vector w_start,
    learning rate eta, and N_epochs epochs.

    Outputs:
        w: A (D, ) shaped array containing the final weight vector.
        losses: A (N_epochs, ) shaped array containing the losses from all iterations.
    """
    #=====
    # TODO: Implement the SGD algorithm.
    #=====

    totalLoss = []
    weights = w_start

    #start loss func (for some reason it won't work when i call the function?)
    #Python TypeError: 'list' object is not callable.
    #and
    #TypeError: 'numpy.float64' object is not callable
    predict = []

    for x in X:
        predict.append(np.inner(weights, x))
    predict = np.asarray(predict)

    loss = 0
    for i in range(len(predict)):
        loss += (predict[i] - Y[i]) ** 2
    #end loss func

    #totalLoss.append(loss)

    for n in range(N_epochs):
        for i in range(len(X)):
            g = gradient(X[i], Y[i], weights)
            weights -= eta * g

        #start loss func
        predict = []
        for x in X:
            predict.append(np.inner(weights, x))
        predict = np.asarray(predict)

        loss = 0
        assert(len(predict) == len(Y))
        for i in range(len(predict)):
            loss += (predict[i] - Y[i]) ** 2

        #currLoss = loss(X, Y, weights)
        totalLoss.append(loss)
        #print(str(n) + ": " + str(loss) + ", " + str(weights))

    return np.asarray(weights), np.asarray(totalLoss)


```

Problem 4D: Visualization

Dataset


We'll start off by generating two simple 2-dimensional datasets. For simplicity we do not consider separate training and test sets.

```
In [83]: X1, Y1 = generate_dataset1()
         plot_dataset(X1, Y1)
```



```
Out[83]: (<Figure size 432x288 with 2 Axes>,
         <matplotlib.axes._subplots.AxesSubplot at 0x11708fa90>)
```

```
In [84]: X2, Y2 = generate_dataset2()
         plot_dataset(X2, Y2)
```



```
Out[84]: (<Figure size 432x288 with 2 Axes>,
         <matplotlib.axes._subplots.AxesSubplot at 0x1170e0750>)
```

SGD from a single point



First, let's visualize SGD from a single starting point:

```
In [85]: # Parameters to feed the SGD.  
# <FR> changes the animation speed.  
params = ({'w_start': [0.01, 0.01], 'eta': 0.00001},)  
N_epochs = 1000  
FR = 20  
  
# Let's animate it!  
anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR)  
HTML(anim.to_html5_video())
```

Performing SGD with parameters {'w_start': [0.01, 0.01], 'eta': 1e-05} ...

Animating...

Out[85]:




Let's view how the weights change as the algorithm converges:

```
In [86]: # Parameters to feed the SGD.
params = ({'w_start': [0.01, 0.01], 'eta': 0.00001},)
N_epochs = 1000
FR = 20

# Let's do it!
W, _ = SGD(X1, Y1, params[0]['w_start'], params[0]['eta'], N_epochs)
anim = animate_convergence(X1, Y1, W, FR)
HTML(anim.to_html5_video())
```

Animating...

Out[86]:



SGD from multiple points

Now, let's visualize SGD from multiple arbitrary starting points:



```
In [87]: # Parameters to feed the SGD.
# Here, we specify each different set of initializations as a dictionary.
params = (
    {'w_start': [-0.8, -0.3], 'eta': 0.00001},
    {'w_start': [-0.9, 0.4], 'eta': 0.00001},
    {'w_start': [-0.4, 0.9], 'eta': 0.00001},
    {'w_start': [0.8, 0.8], 'eta': 0.00001},
)
N_epochs = 1000
FR = 20

# Let's go!
anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR)
HTML(anim.to_html5_video())
```

Performing SGD with parameters {'w_start': [-0.8, -0.3], 'eta': 1e-05} ...
 Performing SGD with parameters {'w_start': [-0.9, 0.4], 'eta': 1e-05} ...
 Performing SGD with parameters {'w_start': [-0.4, 0.9], 'eta': 1e-05} ...
 Performing SGD with parameters {'w_start': [0.8, 0.8], 'eta': 1e-05} ...

Animating...

Out[87]:



Let's do the same thing but with a different dataset:



```
In [88]: # Parameters to feed the SGD.
params = (
    {'w_start': [-0.8, -0.3], 'eta': 0.00001},
    {'w_start': [-0.9, 0.4], 'eta': 0.00001},
    {'w_start': [-0.4, 0.9], 'eta': 0.00001},
    {'w_start': [0.8, 0.8], 'eta': 0.00001},
)
N_epochs = 1000
FR = 20

# Animate!
anim = animate_sgd_suite(SGD, loss, X2, Y2, params, N_epochs, FR)
HTML(anim.to_html5_video())

Performing SGD with parameters {'w_start': [-0.8, -0.3], 'eta': 1e-05} ...
Performing SGD with parameters {'w_start': [-0.9, 0.4], 'eta': 1e-05} ...
Performing SGD with parameters {'w_start': [-0.4, 0.9], 'eta': 1e-05} ...
Performing SGD with parameters {'w_start': [0.8, 0.8], 'eta': 1e-05} ...

Animating...
```

Out[88]:



Problem 4E: SGD with different step sizes

Now, let's visualize SGD with different step sizes (eta):

(For ease of visualization: the trajectories are ordered from left to right by increasing eta value. Also, note that we use smaller values of N_epochs and FR here for easier visualization.)



```
In [89]: # Parameters to feed the SGD.
params = (
    {'w_start': [0.7, 0.8], 'eta': 0.00001},
    {'w_start': [0.2, 0.8], 'eta': 0.00005},
    {'w_start': [-0.2, 0.7], 'eta': 0.0001},
    {'w_start': [-0.6, 0.6], 'eta': 0.0002},
)
N_epochs = 100
FR = 2

# Go!
anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR, ms=2)
HTML(anim.to_html5_video())

Performing SGD with parameters {'w_start': [0.7, 0.8], 'eta': 1e-05} ...
Performing SGD with parameters {'w_start': [0.2, 0.8], 'eta': 5e-05} ...
Performing SGD with parameters {'w_start': [-0.2, 0.7], 'eta': 0.0001} ...
Performing SGD with parameters {'w_start': [-0.6, 0.6], 'eta': 0.0002} ...

Animating...
```

Out[89]:



Plotting SGD Convergence

Let's visualize the difference in convergence rates a different way. Plot the loss with respect to epoch (iteration) number for each value of eta on the same graph.


```
In [90]: '''Plotting SGD convergence'''

#=====
# TODO: For the given learning rates, plot the
# loss for each epoch.
#=====

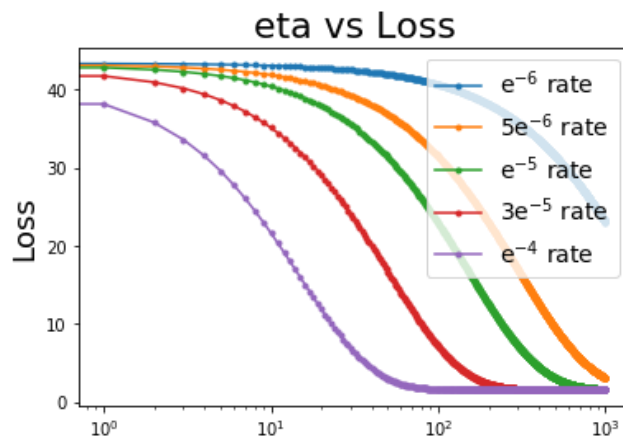
eta_vals = [1e-6, 5e-6, 1e-5, 3e-5, 1e-4]
w_start = [0.01, 0.01]
N_epochs = 1000

weights = []
loss = []

for eta in eta_vals:
    initial = [0.001, 0.001, 0.001, 0.001, 0.001]
    finalWeights, allLoss = SGD(X1, Y1, w_start, eta, N_epochs)
    weights.append(finalWeights)
    loss.append(allLoss)

fig = plt.figure()
x = epochs
plt.title('eta vs Loss', fontsize = 22)
plt.plot(loss[0], marker = '.')
plt.plot(loss[1], marker = '.')
plt.plot(loss[2], marker = '.')
plt.plot(loss[3], marker = '.')
plt.plot(loss[4], marker = '.')
plt.legend(('e-6 rate', '5e-6 rate', 'e-5 rate', '3e-5 rate',
'e-4 rate'), loc = 'upper right', fontsize = 14)
plt.xscale('log')
plt.ylabel('Loss', fontsize = 18)
```

Out[90]: Text(0, 0.5, 'Loss')



Clearly, a big step size results in fast convergence! Why don't we just set eta to a really big value, then? Say, eta=1?

(Again, note that the FR is lower for this animation.)

```
In [91]: # Parameters to feed the SGD.
params = ({'w_start': [0.01, 0.01], 'eta': 1},)
N_epochs = 100
FR = 2

# Voila!
anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR, ms=2)
HTML(anim.to_html5_video())

Performing SGD with parameters {'w_start': [0.01, 0.01], 'eta': 1} ...

-----
TypeError                                Traceback (most recent call last)
<ipython-input-91-6f8ee3a8be68> in <module>
      5
      6 # Voila!
----> 7 anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR, ms=2)
      8 HTML(anim.to_html5_video())

~/Downloads/CS155_SET1/sgd_helper.py in animate_sgd_suite(SGD, loss, X, Y, param
s, N_epochs, FR, ms)
    121
    122     # Get the loss grid and plot it.
--> 123     w_grid, loss_grid = get_loss_grid((-1, 1, 100), (-1, 1, 100), X, Y,
loss)
    124     fig, ax = plot_loss_function(w_grid[0], w_grid[1], loss_grid)
    125

~/Downloads/CS155_SET1/sgd_helper.py in get_loss_grid(x_params, y_params, X, Y,
loss)
    77         for j in range(len(loss_grid[0])):
    78             w = np.array([w_grid[0][i, j], w_grid[1][i, j]])
----> 79             loss_grid[i, j] = loss(X, Y, w)
    80
    81     return w_grid, loss_grid

TypeError: 'list' object is not callable
```

Just for fun, let's try eta=10 as well. What happens? (Hint: look at W)

```
In [ ]: # Parameters to feed the SGD.
w_start = [0.01, 0.01]
eta = 10
N_epochs = 100

# Presto!
W, losses = SGD(X1, Y1, w_start, eta, N_epochs)
```

Extra Visualization (not part of the homework problem)

One final visualization! What happens if the loss function has multiple optima?

```
In [ ]: # Import different SGD & loss functions.
        # In particular, the loss function has multiple optima.
        from sgd_multiopt_helper import SGD, loss

        # Parameters to feed the SGD.
        params = (
            {'w_start': [0.9, 0.9], 'eta': 0.01},
            {'w_start': [0.0, 0.0], 'eta': 0.01},
            {'w_start': [-0.8, 0.6], 'eta': 0.01},
            {'w_start': [-0.8, -0.6], 'eta': 0.01},
            {'w_start': [-0.4, -0.3], 'eta': 0.01},
        )
        N_epochs = 100
        FR = 2

        # One more time!
        anim = animate_sgd_suite(SGD, loss, X1, Y1, params, N_epochs, FR, ms=2)
        HTML(anim.to_html5_video())
```

In []:

In []:

In []:

In []:

In []:

Problem 4, Parts F-H: Stochastic Gradient Descent with a Larger Dataset

Use this notebook to write your code for problem 4 parts F-H by filling in the sections marked `# TODO` and running all cells.

```
In [3]: # Setup.  
import csv  
import numpy as np  
import matplotlib.pyplot as plt  
import math  
import random  
  
%matplotlib inline
```

Problem 4F: Perform SGD with the new dataset

For the functions below, you may re-use your code from parts 4C-E. Note that you can now modify your SGD function to return the final weight vector instead of the weights after every epoch.

```

In [4]: def loss(X, Y, w):
        '''
        Calculate the squared loss function.

        Inputs:
            X: A (N, D) shaped numpy array containing the data points.
            Y: A (N, ) shaped numpy array containing the (float) labels of the data points.
            w: A (D, ) shaped numpy array containing the weight vector.

        Outputs:
            The loss evaluated with respect to X, Y, and w.
        '''
        predict = []
        for x in X:
            predict.append(np.inner(w, x))
        predict = np.asarray(predict)

        loss = 0
        for i in range(len(predict)):
            loss += (predict[i] - Y[i]) ** 2

        return loss

def gradient(x, y, w):
    '''
    Calculate the gradient of the loss function with respect to
    a single point (x, y), and using weight vector w.

    Inputs:
        x: A (D, ) shaped numpy array containing a single data point.
        y: The float label for the data point.
        w: A (D, ) shaped numpy array containing the weight vector.

    Output:
        The gradient of the loss with respect to x, y, and w.
    '''

    #=====
    # TODO: Implement the gradient of the loss function.
    #=====
    grad = -2 * (y - np.inner(w, x)) * x
    return grad

```

```

In [ ]: def SGD(X, Y, w_start, eta, N_epochs):
        '''
        Perform SGD using dataset (X, Y), initial weight vector w_start,
        learning rate eta, and N_epochs epochs.

        Outputs:
            w: A (D, ) shaped array containing the final weight vector.
            losses: A (N_epochs, ) shaped array containing the losses from all iterations.
        '''

        #=====
        # TODO: Implement the SGD algorithm.
        #=====

        totalLoss = []
        weights = w_start

        #start loss func (for some reason it won't work when i call the function?)
        #Python TypeError: 'list' object is not callable.
        #and
        #TypeError: 'numpy.float64' object is not callable
        predict = []
        for x in X:
            predict.append(np.inner(weights, x))
        predict = np.asarray(predict)

        loss = 0
        assert(len(predict) == len(Y))
        for i in range(len(predict)):
            loss += (predict[i] - Y[i]) ** 2
        #end loss func

        totalLoss.append(loss)

        for n in range(N_epochs):
            assert(len(X) == 1000)

            for i in range(len(X)):
                g = gradient(X[i], Y[i], weights)
                assert(len(g) == len(weights))
                weights -= eta * g

            #start loss func
            predict = []
            for x in X:
                predict.append(np.inner(weights, x))
            predict = np.asarray(predict)

            loss = 0
            assert(len(predict) == len(Y))
            for i in range(len(predict)):
                loss += (predict[i] - Y[i]) ** 2

            #currLoss = loss(X, Y, weights)
            totalLoss.append(loss)

        return np.asarray(weights), np.asarray(totalLoss)

```

Next, we need to load the dataset. In doing so, the following function may be helpful:

```
In [5]: def load_data(filename):
        """
        Function loads data stored in the file filename and returns it as a numpy ndarray.

        Inputs:
            filename: GeneratorExitiven as a string.

        Outputs:
            Data contained in the file, returned as a numpy ndarray
        """
        return np.loadtxt(filename, skiprows=1, delimiter=',')

def getData(data):
    """
    This function takes the raw data and returns two arrays for inputs(x)
    and outputs(y)
    """
    x = []
    y = []
    arr = np.asarray([1.0])
    for i in data:
        x.append(np.concatenate((arr, i[:-1]), axis = 0))
        y.append(i[(len(i) - 1)])
    return np.asarray(x), np.asarray(y)
```

Now, load the dataset in `sgd_data.csv` and run SGD using the given parameters; print out the final weights.

```
In [6]: #=====
        # TODO:
        # (1) load the dataset
        # (2) run SGD using the given parameters
        # (3) print out the final weights.
        #=====
        weights = []
        loss = []
        numEpochs = 800
        epochs = [800, 900, 1000, 1100, 1200, 1300]
        steps = [math.exp(-10), math.exp(-11), math.exp(-12), math.exp(-13), math.exp(-14), math.exp(-15)]
        stp = math.exp(-15)
        allData = load_data("data/sgd_data.csv")
        x, y = getData(allData)

        initial = [0.001, 0.001, 0.001, 0.001, 0.001]
        # X, Y, w_start, eta, N_epochs
        fw, _ = SGD(x, y, initial, stp, numEpochs)
        print(fw)

        for e in epochs:
            for step in steps:
                initial = [0.001, 0.001, 0.001, 0.001, 0.001]
                fw, totalLoss = SGD(x, y, initial, step, numEpochs)
                weights.append(fw)
                loss.append(totalLoss)

        [ -0.22720591  -5.94229011   3.94369494 -11.72402388   8.78549375]
```

Problem 4G: Convergence of SGD


This problem examines the convergence of SGD for different learning rates. Please implement your code in the cell below:


```
In [10]: #=====
# TODO: create a plot showing the convergence
# of SGD for the different learning rates.
#=====
print(loss[0])

fig = plt.figure()
x = epochs
plt.title('Training Error vs. Epochs', fontsize = 22)
plt.plot(loss[0], marker = '.')
plt.plot(loss[1], marker = '.')
plt.plot(loss[2], marker = '.')
plt.plot(loss[3], marker = '.')
plt.plot(loss[4], marker = '.')
plt.plot(loss[5], marker = '.')
plt.legend(('e-10 rate', 'e-11 rate', 'e-12 rate', 'e-13 rate',
            'e-14 rate', 'e-15 rate'), loc = 'upper right', fontsize = 14)
plt.xlabel('Number of Epochs (log scale)', fontsize = 18)
plt.xscale('log')
plt.ylabel('Training Error', fontsize = 18)
```

[2874307.79361655	500820.74432231	116012.3436524	29495.32628487
9933.17467149	5483.21974217	4463.39592301	4227.13315194
4171.40613985	4157.83834506	4154.34642156	4153.36282484
4153.04781185	4152.93004034	4152.87821425	4152.85132614
4152.83483885	4152.82299375	4152.81332182	4152.80471177
4152.79664943	4152.78889313	4152.78132811	4152.77389936
4152.76658012	4152.75935713	4152.75222348	4152.74517535
4152.73821038	4152.73132692	4152.72452366	4152.71779947
4152.71115329	4152.70458412	4152.698091	4152.69167298
4152.68532912	4152.6790585	4152.67286023	4152.6667334
4152.66067713	4152.65469055	4152.64877281	4152.64292305
4152.63714043	4152.63142413	4152.62577334	4152.62018724
4152.61466504	4152.60920596	4152.60380922	4152.59847405
4152.59319971	4152.58798544	4152.58283051	4152.57773419
4152.57269577	4152.56771453	4152.56278978	4152.55792082
4152.55310699	4152.54834759	4152.54364198	4152.53898949
4152.53438948	4152.5298413	4152.52534434	4152.52089795
4152.51650154	4152.51215449	4152.50785621	4152.5036061
4152.49940357	4152.49524806	4152.491139	4152.48707581
4152.48305795	4152.47908487	4152.47515603	4152.47127089
4152.46742893	4152.46362962	4152.45987246	4152.45615693
4152.45248254	4152.44884879	4152.4452552	4152.44170127
4152.43818653	4152.43471053	4152.43127278	4152.42787283
4152.42451024	4152.42118454	4152.41789531	4152.41464211
4152.41142451	4152.40824207	4152.4050944	4152.40198106
4152.39890165	4152.39585578	4152.39284303	4152.38986302
4152.38691535	4152.38399965	4152.38111554	4152.37826264
4152.37544058	4152.372649	4152.36988754	4152.36715584
4152.36445355	4152.36178033	4152.35913583	4152.35651971
4152.35393165	4152.3513713	4152.34883835	4152.34633248
4152.34385336	4152.34140069	4152.33897415	4152.33657343
4152.33419825	4152.33184829	4152.32952326	4152.32722288
4152.32494686	4152.3226949	4152.32046674	4152.31826209
4152.31608069	4152.31392226	4152.31178653	4152.30967325
4152.30758215	4152.30551298	4152.30346548	4152.30143941
4152.29943451	4152.29745054	4152.29548727	4152.29354444
4152.29162184	4152.28971922	4152.28783635	4152.28597301
4152.28412898	4152.28230403	4152.28049795	4152.27871052
4152.27694152	4152.27519075	4152.273458	4152.27174307
4152.27004574	4152.26836582	4152.26670312	4152.26505743
4152.26342857	4152.26181634	4152.26022055	4152.25864103
4152.25707758	4152.25553002	4152.25399818	4152.25248187
4152.25098093	4152.24949517	4152.24802444	4152.24656855
4152.24512735	4152.24370067	4152.24228834	4152.24089021
4152.23950611	4152.2381359	4152.2367794	4152.23543649
4152.23410699	4152.23279076	4152.23148766	4152.23019754
4152.22892025	4152.22765565	4152.2264036	4152.22516397
4152.22393661	4152.22272139	4152.22151818	4152.22032684
4152.21914724	4152.21797926	4152.21682277	4152.21567763
4152.21454373	4152.21342094	4152.21230914	4152.21120821
4152.21011804	4152.20903849	4152.20796947	4152.20691085
4152.20586252	4152.20482437	4152.20379629	4152.20277817
4152.2017699	4152.20077137	4152.19978248	4152.19880313
4152.19783321	4152.19687262	4152.19592126	4152.19497903
4152.19404583	4152.19312157	4152.19220615	4152.19129947
4152.19040144	4152.18951197	4152.18863097	4152.18775835
4152.18689402	4152.18603788	4152.18518986	4152.18434986
4152.1835178	4152.18269361	4152.18187718	4152.18106845
4152.18026732	4152.17947373	4152.17868758	4152.17790881
4152.17713733	4152.17637307	4152.17561595	4152.1748659
4152.17412284	4152.1733867	4152.1726574	4152.17193488
4152.17121907	4152.17050989	4152.16980728	4152.16911117
4152.16842149	4152.16773817	4152.16706115	4152.16639036
4152.16572574	4152.16506722	4152.16441474	4152.16376825
4152.16312767	4152.16249294	4152.16186402	4152.16124083

Out[10]: Text(0, 0.5, 'Training Error')



Problem 4H

Provide your code for computing the least-squares analytical solution below.

```
In [9]: #=====
# TODO: implement the least-squares
# analytical solution.
#=====

x, y = getData(allData)
x = np.matrix(x)
y = np.matrix(y)
ret = np.linalg.inv(x.transpose() * x) * x.transpose() * y.transpose()

print(list(ret))

[matrix([[ -0.31644251]]), matrix([[ -5.99157048]]), matrix([[ 4.01509955]]), matri
x([[ -11.93325972]]), matrix([[ 8.99061096]])]
```

In []: