

Problem 3

Use this notebook to write your code for problem 3.

```
In [4]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

3D - Convolutional network

As in problem 2, we have conveniently provided for your use code that loads and preprocesses the MNIST data.

```
In [5]: # load MNIST data into PyTorch format
import torch
import torchvision
import torchvision.transforms as transforms

# set batch size
batch_size = 32

# load training data downloaded into data/ folder
mnist_training_data = torchvision.datasets.MNIST('data/', train=True, download=True,
                                                transform=transforms.ToTensor())
# transforms.ToTensor() converts batch of images to 4-D tensor and normalizes 0-255 to 0-1.0
training_data_loader = torch.utils.data.DataLoader(mnist_training_data,
                                                    batch_size=batch_size,
                                                    shuffle=True)

# load test data
mnist_test_data = torchvision.datasets.MNIST('data/', train=False, download=True,
                                              transform=transforms.ToTensor())
test_data_loader = torch.utils.data.DataLoader(mnist_test_data,
                                                batch_size=batch_size,
                                                shuffle=False)
```

```
In [6]: # look at the number of batches per epoch for training and validation
print(f'{len(training_data_loader)} training batches')
print(f'{len(training_data_loader) * batch_size} training samples')
print(f'{len(test_data_loader)} validation batches')
```

```
1875 training batches
60000 training samples
313 validation batches
```

```
In [65]: # sample model
import torch.nn as nn

class ConvNet(nn.Module):
    def __init__(self):
        super(ConvNet, self).__init__()
        self.first = nn.Sequential(
            nn.Conv2d(1, 16, kernel_size=5, stride=1, padding=2),
            nn.BatchNorm2d(16),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.second = nn.Sequential(
            nn.Conv2d(16, 32, kernel_size=5, stride=1, padding=2),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))
        self.Dropout = nn.Dropout()
        # 7 * 7 * 32
        self.fc1 = nn.Linear(7*7*32, 600)
        self.fc2 = nn.Linear(600, 10)

    def forward(self, data):
        out = self.first(data)
        out = self.second(out)
        # reshape
        out = out.reshape(out.size(0), -1)
        out = self.Dropout(out)
        out = self.fc1(out)
        out = self.fc2(out)
        return out

model = ConvNet()
```

```
In [66]: # why don't we take a look at the shape of the weights for each layer
for p in model.parameters():
    print(p.data.shape)
```

```
torch.Size([16, 1, 5, 5])
torch.Size([16])
torch.Size([16])
torch.Size([16])
torch.Size([32, 16, 5, 5])
torch.Size([32])
torch.Size([32])
torch.Size([32])
torch.Size([600, 1568])
torch.Size([600])
torch.Size([10, 600])
torch.Size([10])
```

```
In [67]: # our model has some # of parameters:
count = 0
for p in model.parameters():
    n_params = np.prod(list(p.data.shape)).item()
    count += n_params
print(f'total params: {count}')
```

```
total params: 960754
```

```
In [68]: # For a multi-class classification problem  
#import torch.optim as optim  
criterion = nn.CrossEntropyLoss()  
optimizer = optim.Adam(model.parameters(), lr=.001)
```

```

In [69]: # Train the model for 10 epochs, iterating on the data in batches
n_epochs = 10

# store metrics
training_accuracy_history = np.zeros([n_epochs, 1])
training_loss_history = np.zeros([n_epochs, 1])
validation_accuracy_history = np.zeros([n_epochs, 1])
validation_loss_history = np.zeros([n_epochs, 1])

for epoch in range(n_epochs):
    print(f'Epoch {epoch+1}/10:', end='')
    train_total = 0
    train_correct = 0
    # train
    model.train()
    for i, data in enumerate(training_data_loader):
        images, labels = data
        optimizer.zero_grad()
        # forward pass
        output = model(images)
        # calculate categorical cross entropy loss
        loss = criterion(output, labels)
        # backward pass
        loss.backward()
        optimizer.step()

        # track training accuracy
        _, predicted = torch.max(output.data, 1)
        train_total += labels.size(0)
        train_correct += (predicted == labels).sum().item()
        # track training loss
        training_loss_history[epoch] += loss.item()
        # progress update after 180 batches (~1/10 epoch for batch size 32)
        if i % 180 == 0: print('.', end='')
    training_loss_history[epoch] /= len(training_data_loader)
    training_accuracy_history[epoch] = train_correct / train_total
    print(f'\n\tloss: {training_loss_history[epoch,0]:0.4f}, acc: {training_accuracy_history[epoch,0]:0.4f}', end='')

    # validate
    test_total = 0
    test_correct = 0
    with torch.no_grad():
        model.eval()
        for i, data in enumerate(test_data_loader):
            images, labels = data
            # forward pass
            output = model(images)
            # find accuracy
            _, predicted = torch.max(output.data, 1)
            test_total += labels.size(0)
            test_correct += (predicted == labels).sum().item()
            # find loss
            loss = criterion(output, labels)
            validation_loss_history[epoch] += loss.item()
        validation_loss_history[epoch] /= len(test_data_loader)
        validation_accuracy_history[epoch] = test_correct / test_total
    print(f', val loss: {validation_loss_history[epoch,0]:0.4f}, val acc: {validation_accuracy_history[epoch,0]:0.4f}')

```

```
Epoch 1/10:.....
    loss: 0.2103, acc: 0.9369, val loss: 0.0465, val acc: 0.9843
Epoch 2/10:.....
    loss: 0.0968, acc: 0.9709, val loss: 0.0462, val acc: 0.9848
Epoch 3/10:.....
    loss: 0.0722, acc: 0.9785, val loss: 0.0308, val acc: 0.9896
Epoch 4/10:.....
    loss: 0.0634, acc: 0.9807, val loss: 0.0355, val acc: 0.9889
Epoch 5/10:.....
    loss: 0.0582, acc: 0.9825, val loss: 0.0327, val acc: 0.9886
Epoch 6/10:.....
    loss: 0.0504, acc: 0.9847, val loss: 0.0305, val acc: 0.9905
Epoch 7/10:.....
    loss: 0.0477, acc: 0.9854, val loss: 0.0318, val acc: 0.9896
Epoch 8/10:.....
    loss: 0.0441, acc: 0.9859, val loss: 0.0301, val acc: 0.9909
Epoch 9/10:.....
    loss: 0.0431, acc: 0.9870, val loss: 0.0314, val acc: 0.9902
Epoch 10/10:.....
    loss: 0.0390, acc: 0.9881, val loss: 0.0232, val acc: 0.9925
```

Above, we output the training loss/accuracy as well as the validation loss and accuracy. Not bad! Let's see if you can do better.