

Problem 3

Authors: Sid Murching, Suraj Nair, Alex Cui

```
In [9]: import numpy as np
        from P3CHelpers import *
        from keras.models import Sequential
        import sys
```

3D:

Fill in the generate_traindata and find_most_similar_pairs functions

```

In [10]: def get_word_repr(word_to_index, word):
    """
    Returns one-hot-encoded feature representation of the specified word given
    a dictionary mapping words to their one-hot-encoded index.

    Arguments:
        word_to_index: Dictionary mapping words to their corresponding index
                        in a one-hot-encoded representation of our corpus.

        word:          String containing word whose feature representation we wish
                        to compute.

    Returns:
        feature_representation: Feature representation of the passed-in word.
    """
    unique_words = word_to_index.keys()
    # Return a vector that's zero everywhere besides the index corresponding to <word>
    feature_representation = np.zeros(len(unique_words))
    feature_representation[word_to_index[word]] = 1
    return feature_representation

def generate_traindata(word_list, word_to_index, window_size=4):
    """
    Generates training data for Skipgram model.

    Arguments:
        word_list:      Sequential list of words (strings).
        word_to_index:  Dictionary mapping words to their corresponding index
                        in a one-hot-encoded representation of our corpus.

        window_size:    Size of Skipgram window.
                        (use the default value when running your code).

    Returns:
        (trainX, trainY): A pair of matrices (trainX, trainY) containing training
                        points (one-hot-encoded vectors representing individual words) and
                        their corresponding labels (also one-hot-encoded vectors representing words).

        For each index i, trainX[i] should correspond to a word in
        <word_list>, and trainY[i] should correspond to one
        of the words within
        a window of size <window_size> of trainX[i].
    """
    trainX = []
    trainY = []
    numWords = len(word_to_index)
    trainX = []
    trainY = []
    allZeroes = [0 for i in range(numWords)]
    for i in range(len(word_list)):
        for j in range(-window_size, window_size + 1):
            if i + j >= 0 and i + j < len(word_list) and j != 0:
                trainXVector = get_word_repr(word_to_index, word_list[i])
                trainX.append(trainXVector)
                trainYVector = get_word_repr(word_to_index, word_list[i+j])
                trainY.append(trainYVector)

    return (np.array(trainX), np.array(trainY))

```

```

In [12]: def find_most_similar_pairs(filename, num_latent_factors):
        """
        Find the most similar pairs from the word embeddings computed from
        a body of text

        Arguments:
            filename: Text file to read and train embeddings from
            num_latent_factors: The number of latent factors / the size of the embedding
        """
        # Load in a list of words from the specified file; remove non-alphanumeric characters
        # and make all chars lowercase.
        sample_text = load_word_list(filename)

        # Create dictionary mapping unique words to their one-hot-encoded index
        word_to_index = generate_onehot_dict(sample_text)
        # Create training data using default window size
        trainX, trainY = generate_traindata(sample_text, word_to_index)

        # TODO: 1) Create and train model in Keras.

        # vocab_size = number of unique words in our text file. Will be useful when adding layers
        # to your neural network
        vocab_size = len(word_to_index)
        model = Sequential()
        model.add(Dense(num_latent_factors, input_dim=(vocab_size)))
        model.add(Dense(vocab_size))
        model.add(Activation('softmax'))
        model.compile(loss='categorical_crossentropy', optimizer='rmsprop',
                      metrics=['accuracy'])
        fit = model.fit(trainX, trainY)

        print("Hidden layer shape" + str(model.layers[0].get_weights()[0].shape))
        print("Output layer shape" + str(model.layers[1].get_weights()[0].shape))
        # TODO: 2) Extract weights for hidden layer, set <weights> variable below

        weights = model.layers[0].get_weights()[0]

        # Find and print most similar pairs
        similar_pairs = most_similar_pairs(weights, word_to_index)
        for pair in similar_pairs[:30]:
            print(pair)

```

In [6]:

In []:

3G:

Run the function below and report your results for dr_seuss.txt.

In []: find_most_similar_pairs('data/dr_seuss.txt', 10)

In []: